

# Information

Matthew Blasa

Student ID: 001781641

MS Data Analytics (05/01/2021)

Program Mentor: Kirk Kelly

(503)805-0297

mblasa@wgu.edu

## A1. Question to Analyze

What is the average amount of data that we estimate the customer will use in the future? Can this be predicted accurately from a list of explanatory variables?

## A2 Analysis Benefits

Stakeholders in the company will benefit by knowing what factors affect the average amount of GB used in a year by the customer. This will provide insight for decisions in whether or not to expand customer data limits, provide unlimited (or metered) media streaming, and expand infrastructure and maintenance support for increased bandwidth demands.

## Regression Methods

### B1. Assumptions of Logistic Regression Model

- Linear relationship between the dependent and independent variables.
- The independent variables are not too highly correlated with each other.
- Observations are selected independently & randomly from the population.
- Residuals should normally be distributed with a mean of zero. The variance of the residuals is constant

### B2. Benefits of Using Python

I will use the Python, since I use it at work to perform EDAs and clean data. I'll be using jupyter notebook, since it allows me to use markdowns to answer questions and program at the same time. To save time, I will be using several data science Python libraries to simplify the problem. The following will be used:

- NumPy - to work with arrays
- Pandas - used to create dataframes.
- Matplotlib - plotting charts
- Scikit-learn - for PCA, Machine Learning, and Normalization.

- SciPy - used for mathematical transformations.
- Seaborn - visualization of more complex graphs

## B3. Justification

Multiple regression is an appropriate technique to analyze the research question because our target variable, predicting a real number of GBs per year, is a continuous variable.

Several explanatory variables help to predict how much data a customer will use in a given year. When adding or removing independent variables from our regression equation, we determine whether or not they have a positive or negative relationship to our target variable.

## C2. Summary Statistics

Summary statistics analysed & discussed below, including the target variable and all predictor variables needed to gather from the dataset to answer the research question.

## C3. Steps to Prepare the Data

This project uses the data that was previously cleaned in the Data Cleaning Assignment. The majority of the data set is used to determine the amount of average amount of data used, in GB, in a year by the customer. I plan to reduce the list of variables not by what might intuitively belong together, but by looking at numerical calculations that suggest which variables are significant to answering the research question. With variables, I aim to get a predictive model to assist the telecommunications company to predict customers and target customers who use the most

The following steps will be taken to prepare the data for analysis:

1. Import data to pandas dataframe
2. Determine variable types and those which may require further investigation
3. Convert binary variables to yes = 1 and no = 0
4. Investigate potential categorical variables using bar charts, then convert categorical values to dummy variables
5. Drop columns that will not be used in regression analysis

My approach will include:

1. Import data to pandas dataframe
2. Determine variable types and those which may require further investigation
3. Convert binary variables to yes = 1 and no = 0
4. Investigate potential categorical variables using bar charts, then convert categorical values to dummy variables
5. Drop columns that will not be used in regression analysis
6. Find outliers that may create or hide statistical significance using histograms.

7. Substitute missing data with meaningful measures of central tendency (mean, median or mode)

Important to the process is the variable of "Bandwidth\_GB\_Year" (the average yearly amount of data used, in GB, per customer) which will be our target variable.

When analyzing the data we will examine the following continuous predictor variables:

- Children
- Income
- Outage\_sec\_perweek
- Email
- Contacts
- Yearly\_equip\_failure
- MonthlyCharge
- Bandwidth\_GB\_Year

We will also examine the following categorical predictor variables:

- Churn: Whether the customer discontinued service within the last month (yes, no)
- Techie: Whether the customer considers themselves technically inclined (based on customer questionnaire when they signed up for services) (yes, no)
- Contract: The contract term of the customer (month-to-month, one year, two year)
- Port\_modem: Whether the customer has a portable modem (yes, no)
- Tablet: Whether the customer owns a tablet such as iPad, Surface, etc. (yes, no)
- InternetService: Customer's internet service provider (DSL, fiber optic, None)
- Phone: Whether the customer has a phone service (yes, no)
- Multiple: Whether the customer has multiple lines (yes, no)
- OnlineSecurity: Whether the customer has an online security add-on (yes, no)
- OnlineBackup: Whether the customer has an online backup add-on (yes, no)
- DeviceProtection: Whether the customer has device protection add-on (yes, no)
- TechSupport: Whether the customer has a technical support add-on (yes, no)
- StreamingTV: Whether the customer has streaming TV (yes, no)
- StreamingMovies: Whether the customer has streaming movies (yes, no)

We will also examine ordinal predictor variables from the survey responses from customers. In the surveys, customers rated eight customer service factors on a scale from 1 to 8.

- Item1: Timely response
- Item2: Timely fixes
- Item3: Timely replacements
- Item4: Reliability
- Item5: Options
- Item6: Respectful response
- Item7: Courteous exchange

Step 1

```
In [1]: # Import necessary libraries
import pandas as pd
import numpy as np
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
from matplotlib.patches import Rectangle
import seaborn as sns
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import cross_val_predict, train_test_split
from sklearn.linear_model import Lasso, LassoCV, Ridge, RidgeCV
import statsmodels.api as sm
import statsmodels.formula.api as smf
from yellowbrick.regressor import AlphaSelection, PredictionError, ResidualsPlot
import warnings
warnings.filterwarnings('ignore') # Ignore warning messages for readability
```

```
In [2]: # Read in dataset and view head
df = pd.read_csv('churn_clean.csv')
pd.options.display.max_columns = None
df.head()
```

```
Out[2]:
```

	CaseOrder	Customer_id	Interaction	UID	City	State
0	1	K409198	aa90260b-4141-4a24-8e36-b04ce1f4f77b	e885b299883d4f9fb18e39c75155d990	Point Baker	AK
1	2	S120509	fb76459f-c047-4a9d-8af9-e0f7d4ac2524	f2de8bef964785f41a2959829830fb8a	West Branch	MI
2	3	K191035	344d114c-3736-4be5-98f7-c72c281e2d35	f1784cfa9f6d92ae816197eb175d3c71	Yamhill	OR
3	4	D90850	abfa2b40-2d43-4994-b15a-989b8c79e311	dc8a365077241bb5cd5ccd305136b05e	Del Mar	CA
4	5	K662701	68a861fd-0d20-4e51-a587-8a90407ee574	aabb64a116e83fdc4befc1fbab1663f9	Needville	TX

Step 2

```
In [3]: # View list of columns, data types, and missing values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 50 columns):
```

#	Column	Non-Null Count	Dtype
0	CaseOrder	10000 non-null	int64
1	Customer_id	10000 non-null	object
2	Interaction	10000 non-null	object
3	UID	10000 non-null	object
4	City	10000 non-null	object
5	State	10000 non-null	object
6	County	10000 non-null	object
7	Zip	10000 non-null	int64
8	Lat	10000 non-null	float64
9	Lng	10000 non-null	float64
10	Population	10000 non-null	int64
11	Area	10000 non-null	object
12	TimeZone	10000 non-null	object
13	Job	10000 non-null	object
14	Children	10000 non-null	int64
15	Age	10000 non-null	int64
16	Income	10000 non-null	float64
17	Marital	10000 non-null	object
18	Gender	10000 non-null	object
19	Churn	10000 non-null	object
20	Outage_sec_perweek	10000 non-null	float64
21	Email	10000 non-null	int64
22	Contacts	10000 non-null	int64
23	Yearly_equip_failure	10000 non-null	int64
24	Techie	10000 non-null	object
25	Contract	10000 non-null	object
26	Port_modem	10000 non-null	object
27	Tablet	10000 non-null	object
28	InternetService	10000 non-null	object
29	Phone	10000 non-null	object
30	Multiple	10000 non-null	object
31	OnlineSecurity	10000 non-null	object
32	OnlineBackup	10000 non-null	object
33	DeviceProtection	10000 non-null	object
34	TechSupport	10000 non-null	object
35	StreamingTV	10000 non-null	object
36	StreamingMovies	10000 non-null	object
37	PaperlessBilling	10000 non-null	object
38	PaymentMethod	10000 non-null	object
39	Tenure	10000 non-null	float64
40	MonthlyCharge	10000 non-null	float64
41	Bandwidth_GB_Year	10000 non-null	float64
42	Item1	10000 non-null	int64
43	Item2	10000 non-null	int64
44	Item3	10000 non-null	int64
45	Item4	10000 non-null	int64
46	Item5	10000 non-null	int64
47	Item6	10000 non-null	int64
48	Item7	10000 non-null	int64
49	Item8	10000 non-null	int64

dtypes: float64(7), int64(16), object(27)

memorv usage: 3.8+ MB

In [4]:

```
df.columns
```

Out[4]: Index(['CaseOrder', 'Customer\_id', 'Interaction', 'UID', 'City', 'State', 'County', 'Zip', 'Lat', 'Lng', 'Population', 'Area', 'TimeZone', 'Job', 'Children', 'Age', 'Income', 'Marital', 'Gender', 'Churn', 'Outage\_sec\_perweek', 'Email', 'Contacts', 'Yearly\_equip\_failure', 'Techie', 'Contract', 'Port\_modem', 'Tablet', 'InternetService', 'Phone', 'Multiple', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',

```
'PaperlessBilling', 'PaymentMethod', 'Tenure', 'MonthlyCharge',
'Bandwidth_GB_Year', 'Item1', 'Item2', 'Item3', 'Item4', 'Item5',
'Item6', 'Item7', 'Item8'],
dtype='object')
```


## Results of Cleaning:

There are no missing values. The data set is clean and ready to prepare for analysis

```
In [5]: # Importing the libraries
import missingno as msno

# Visualize missing values as a matrix
msno.matrix(df)
```

Out[5]: <AxesSubplot:>



## Step 3

```
In [6]: # Convert binary variables into yes = 1, no = 0 (ref 1)
cols = ['Marital', 'Gender', 'Churn', 'Techie', 'Port_modem', 'Tablet', 'Inter
        'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
df[cols] = df[cols].replace(to_replace = ['No', 'Yes'], value = [0, 1])
```

```
In [7]: df.rename(columns = {'Item1': 'TimelyResponse',
                             'Item2': 'Fixes',
                             'Item3': 'Replacements',
                             'Item4': 'Reliability',
                             'Item5': 'Options',
                             'Item6': 'Respectfulness',
                             'Item7': 'Courteous',
                             'Item8': 'Listening'},
                  inplace=True)
```


```
In [8]: df['Gender'] = [1 if v == 'Male' else 0 for v in df['Gender']]
df['InternetService'] = [1 if v == 'Fiber Optic' else 0 for v in df['InternetService']]
```

```
In [ ]:
```

## Step 4

```
In [9]: # View bar charts for potential categorical variables to determine number of
figure, axes = plt.subplots(nrows=7, ncols=2, figsize=(20,20))
plt.subplot(7, 2, 1)
sns.countplot(data = df, y = 'Techie')
plt.subplot(7, 2, 2)
sns.countplot(data = df, y = 'Gender')
plt.subplot(7, 2, 3)
sns.countplot(data = df, y = 'Port_modem')
plt.subplot(7, 2, 4)
sns.countplot(data = df, y = 'Tablet')
plt.subplot(7, 2, 5)
sns.countplot(data = df, y = 'Phone')
plt.subplot(7, 2, 6)
sns.countplot(data = df, y = 'Multiple')
plt.subplot(7, 2, 7)
sns.countplot(data = df, y = 'OnlineSecurity')
plt.subplot(7, 2, 8)
sns.countplot(data = df, y = 'OnlineBackup')
plt.subplot(7, 2, 9)
sns.countplot(data = df, y = 'DeviceProtection')
plt.subplot(7, 2, 10)
sns.countplot(data = df, y = 'TechSupport')
plt.subplot(7, 2, 11)
sns.countplot(data = df, y = 'PaperlessBilling')
plt.subplot(7, 2, 12)
sns.countplot(data = df, y = 'PaperlessBilling')
plt.subplot(7, 2, 13)
sns.countplot(data = df, y = 'Churn')

plt.subplot(7, 2, 14).set_visible(False)
figure.tight_layout()
plt.show();
```



- Country appears to currently only have one value, so will be treated as a string. The remaining variables will be converted into dummy variables.

In [10]:

df

Out[10]:

	CaseOrder	Customer_id	Interaction	UID	City
0	1	K409198	aa90260b-4141-4a24-8e36-b04ce1f4f77b	e885b299883d4f9fb18e39c75155d990	Point Baker
1	2	S120509	fb76459f-c047-4a9d-8af9-e0f7d4ac2524	f2de8bef964785f41a2959829830fb8a	West Branch
2	3	K191035	344d114c-3736-4be5-98f7-c72c281e2d35	f1784cfa9f6d92ae816197eb175d3c71	Yamhill



	CaseOrder	Customer_id	Interaction	UID	City
3	4	D90850	abfa2b40-2d43-4994-b15a-989b8c79e311	dc8a365077241bb5cd5ccd305136b05e	Del Mar
4	5	K662701	68a861fd-0d20-4e51-a587-8a90407ee574	aabb64a116e83fdc4befc1fbab1663f9	Needville
...	...	...	...	...	...
9995	9996	M324793	45deb5a2-ae04-4518-bf0b-c82db8dbe4a4	9499fb4de537af195d16d046b79fd20a	Mount Holly
9996	9997	D861732	6e96b921-0c09-4993-bbda-a1ac6411061a	c09a841117fa81b5c8e19afec2760104	Clarksville
9997	9998	I243405	e8307ddf-9a01-4fff-bc59-4742e03fd24f	9c41f212d1e04dca84445019bbc9b41c	Mobeetie
9998	9999	I641617	3775ccfc-0052-4107-81ae-9657f81ecdf3	3e1f269b40c235a1038863ecf6b7a0df	Carrollton
9999	10000	T38070	9de5fb6e-bd33-4995-aec8-f01d0172a499	0ea683a03a3cd544ae8388aab16176	Clarkesville

## Step 5

```
In [11]: # Drop columns not needed for analysis
drops = ['CaseOrder', 'Interaction', 'Customer_id', 'UID', 'County', 'State',
         'Job', 'TimeZone', 'Area', 'Contract', 'InternetService', 'Marital',
         'df2 = df.drop(drops, axis = 1)']
```

```
In [12]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Population                            10000 non-null  int64
1   Children                             10000 non-null  int64
2   Age                                  10000 non-null  int64
3   Income                              10000 non-null  float64
4   Gender                              10000 non-null  int64
5   Churn                               10000 non-null  int64
6   Outage_sec_perweek                   10000 non-null  float64
7   Email                               10000 non-null  int64
8   Contacts                            10000 non-null  int64
9   Yearly_equip_failure                 10000 non-null  int64
10  Techie                              10000 non-null  int64
11  Port_modem                           10000 non-null  int64
12  Tablet                              10000 non-null  int64
13  Phone                               10000 non-null  int64
14  Multiple                             10000 non-null  int64
```

```

15 OnlineSecurity      10000 non-null int64
16 OnlineBackup        10000 non-null int64
17 DeviceProtection    10000 non-null int64
18 TechSupport         10000 non-null int64
19 StreamingTV         10000 non-null int64
20 StreamingMovies     10000 non-null int64
21 PaperlessBilling    10000 non-null int64
22 Tenure              10000 non-null float64
23 MonthlyCharge       10000 non-null float64
24 Bandwidth_GB_Year   10000 non-null float64
25 TimelyResponse      10000 non-null int64
26 Fixes               10000 non-null int64
27 Replacements        10000 non-null int64
28 Reliability         10000 non-null int64
29 Options             10000 non-null int64
30 Respectfulness     10000 non-null int64
31 Courteous           10000 non-null int64
32 Listening            10000 non-null int64
dtypes: float64(5), int64(28)
memory usage: 2.5 MB

```

## VISUALIZATIONS

### Univariate Visualizations

In [13]:

```
df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Population                            10000 non-null  int64
1   Children                             10000 non-null  int64
2   Age                                  10000 non-null  int64
3   Income                               10000 non-null  float64
4   Gender                               10000 non-null  int64
5   Churn                                10000 non-null  int64
6   Outage_sec_perweek                  10000 non-null  float64
7   Email                                10000 non-null  int64
8   Contacts                             10000 non-null  int64
9   Yearly_equip_failure                10000 non-null  int64
10  Techie                               10000 non-null  int64
11  Port_modem                           10000 non-null  int64
12  Tablet                               10000 non-null  int64
13  Phone                                10000 non-null  int64
14  Multiple                             10000 non-null  int64
15  OnlineSecurity                       10000 non-null  int64
16  OnlineBackup                         10000 non-null  int64
17  DeviceProtection                     10000 non-null  int64
18  TechSupport                          10000 non-null  int64
19  StreamingTV                          10000 non-null  int64
20  StreamingMovies                      10000 non-null  int64
21  PaperlessBilling                     10000 non-null  int64
22  Tenure                               10000 non-null  float64
23  MonthlyCharge                        10000 non-null  float64
24  Bandwidth_GB_Year                    10000 non-null  float64
25  TimelyResponse                       10000 non-null  int64
26  Fixes                                10000 non-null  int64
27  Replacements                         10000 non-null  int64
28  Reliability                          10000 non-null  int64
29  Options                              10000 non-null  int64
30  Respectfulness                       10000 non-null  int64
31  Courteous                            10000 non-null  int64

```

```

32 Listening 10000 non-null int64
dtypes: float64(5), int64(28)
memory usage: 2.5 MB

```

```
In [14]: df2.columns
```

```


Out[14]: Index(['Population', 'Children', 'Age', 'Income', 'Gender', 'Churn',
               'Outage_sec_perweek', 'Email', 'Contacts', 'Yearly equip_failure',
               'Techie', 'Port_modem', 'Tablet', 'Phone', 'Multiple', 'OnlineSecurity',
               'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
               'StreamingMovies', 'PaperlessBilling', 'Tenure', 'MonthlyCharge',
               'Bandwidth_GB_Year', 'TimelyResponse', 'Fixes', 'Replacements',
               'Reliability', 'Options', 'Respectfulness', 'Courteous', 'Listening'],
              dtype='object')



```



```


In [15]: # Display boxplots of numeric columns (ref 2)
cols = ['Age', 'Income', 'Outage_sec_perweek', 'Email', 'Contacts', 'Yearly_e
for v in cols:
    sns.boxplot(v, data = df2)
    plt.show()
    print('\n')
    print('\n')

```









```
In [16]: df2.hist(figsize=(20, 15))
```

```
Out[16]: array([[<AxesSubplot:title={'center':'Population'}>,
<AxesSubplot:title={'center':'Children'}>,
<AxesSubplot:title={'center':'Age'}>,
<AxesSubplot:title={'center':'Income'}>,
<AxesSubplot:title={'center':'Gender'}>,
<AxesSubplot:title={'center':'Churn'}>],
[<AxesSubplot:title={'center':'Outage_sec_perweek'}>,
<AxesSubplot:title={'center':'Email'}>,
<AxesSubplot:title={'center':'Contacts'}>,
<AxesSubplot:title={'center':'Yearly equip_failure'}>,
<AxesSubplot:title={'center':'Techie'}>,
<AxesSubplot:title={'center':'Port_modem'}>],
[<AxesSubplot:title={'center':'Tablet'}>,
<AxesSubplot:title={'center':'Phone'}>,
<AxesSubplot:title={'center':'Multiple'}>,
<AxesSubplot:title={'center':'OnlineSecurity'}>,
<AxesSubplot:title={'center':'OnlineBackup'}>,
<AxesSubplot:title={'center':'DeviceProtection'}>],
```

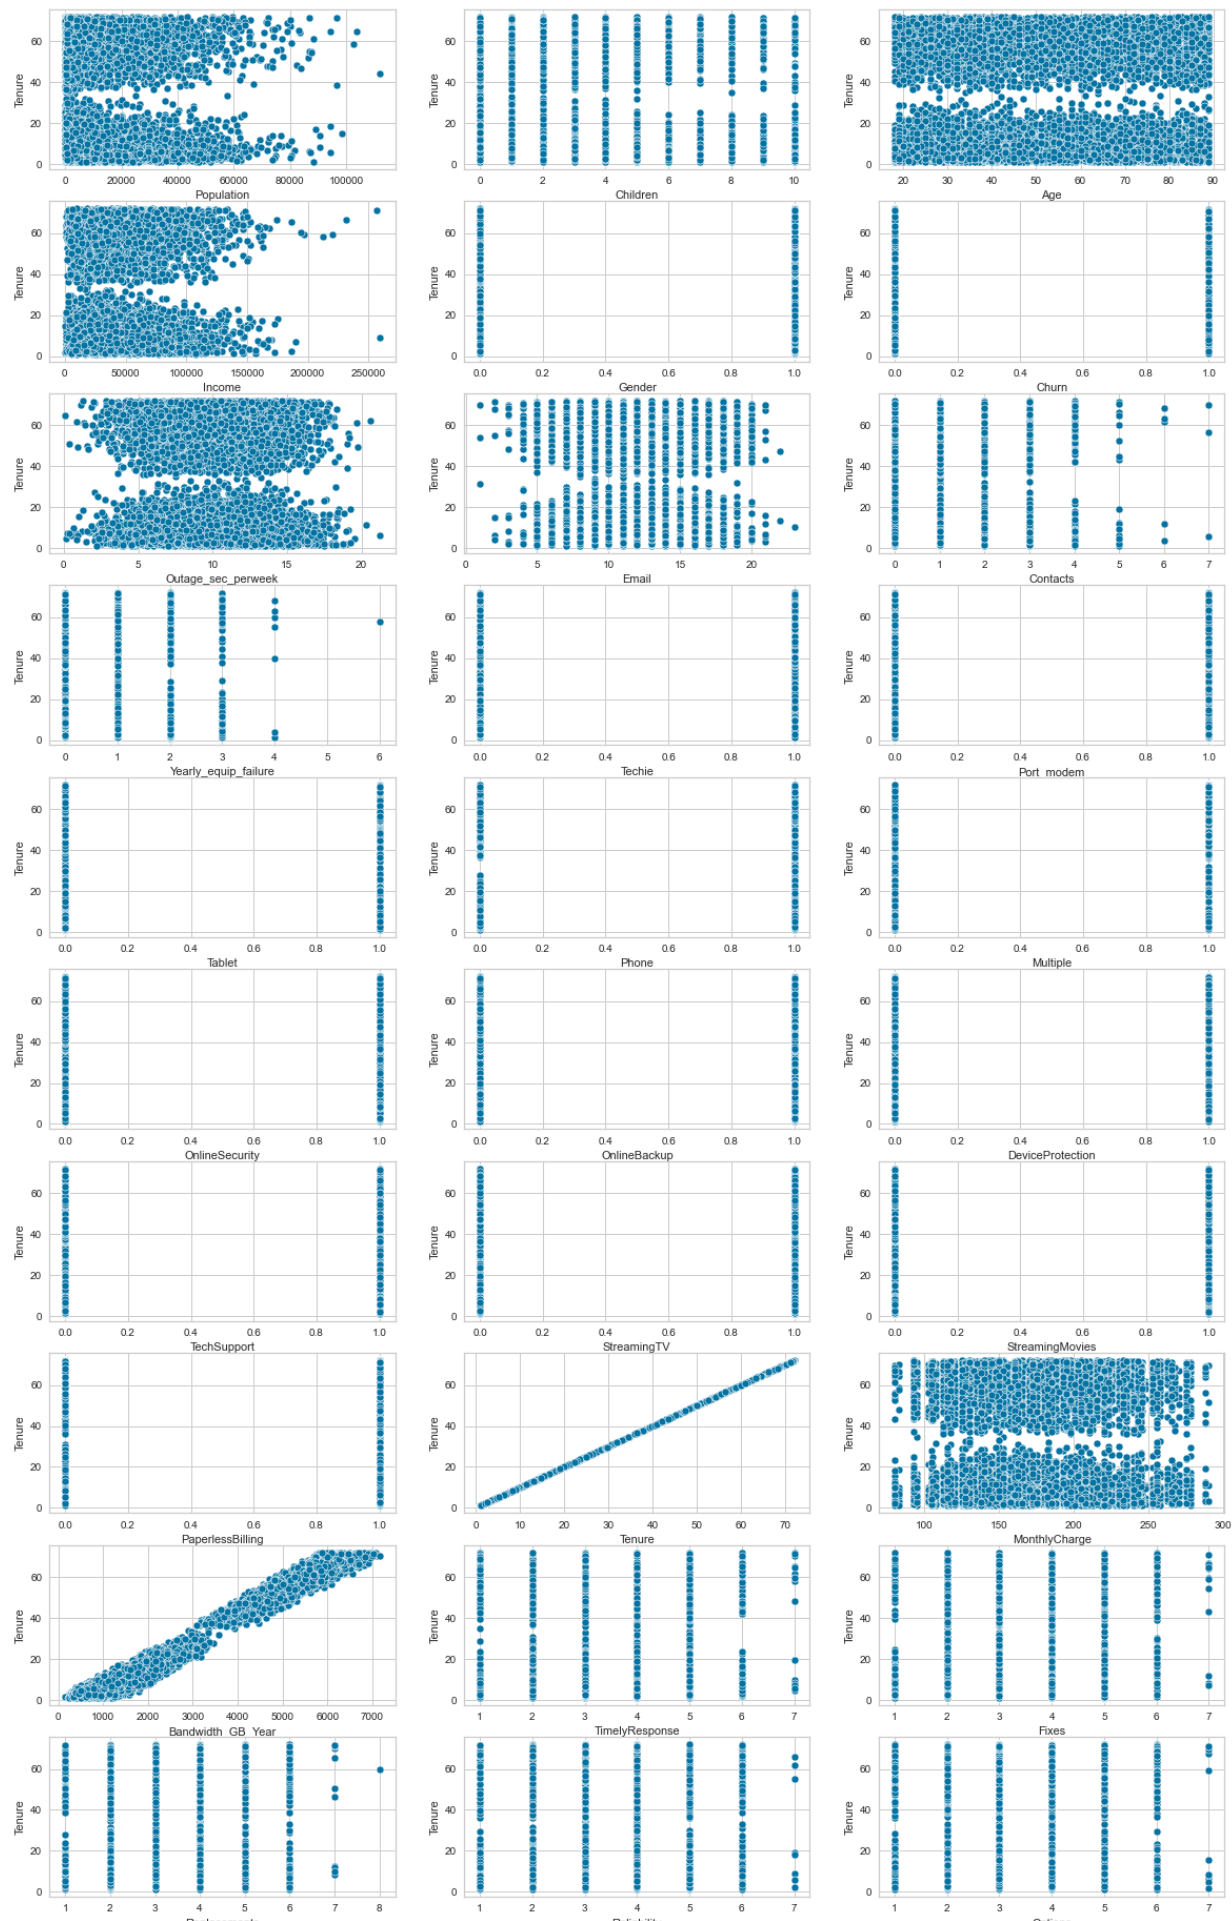
```
[<AxesSubplot:title={'center':'TechSupport'}>,
 <AxesSubplot:title={'center':'StreamingTV'}>,
 <AxesSubplot:title={'center':'StreamingMovies'}>,
 <AxesSubplot:title={'center':'PaperlessBilling'}>,
 <AxesSubplot:title={'center':'Tenure'}>,
 <AxesSubplot:title={'center':'MonthlyCharge'}>],
 [<AxesSubplot:title={'center':'Bandwidth_GB_Year'}>,
 <AxesSubplot:title={'center':'TimelyResponse'}>,
 <AxesSubplot:title={'center':'Fixes'}>,
 <AxesSubplot:title={'center':'Replacements'}>,
 <AxesSubplot:title={'center':'Reliability'}>,
 <AxesSubplot:title={'center':'Options'}>],
 [<AxesSubplot:title={'center':'Respectfulness'}>,
 <AxesSubplot:title={'center':'Courteous'}>,
 <AxesSubplot:title={'center':'Listening'}>], <AxesSubplot:>,
 <AxesSubplot:>], dtype=object)
```



## Bivariate Visualizations

```
In [17]: # Plot scatter plots for each predictive variable on the x axis and Tenure on
count=1
plt.subplots(figsize=(20, 80))
for i in df2.columns:
    plt.subplot(24,3,count)
    sns.scatterplot(df2[i], df2["Tenure"])
    count+=1


plt.show()
```





In [18]:

```
# Display heatmap to view correlation between variables highlighting the Tenu.
t = df2.columns.get_loc('Bandwidth_GB_Year')
l = len(df2.columns)
fig, ax = plt.subplots(figsize=(30,30))
sns.heatmap(df2.corr(), cmap='coolwarm')
ax.add_patch(Rectangle((t, 0), 1, 1, fill=False, edgecolor='blue', lw=3))
ax.add_patch(Rectangle((0, t), 1, 1, fill=False, edgecolor='blue', lw=3))
plt.show()
```



## Prepared Data Set

```
In [19]: # Save cleaned dataframe to CSV
df.to_csv('churn_clean_data_final.csv', index = False, encoding = 'utf-8')
```

## D1. Initial Model

### Initial Model

```
In [20]: # Set up input matrix and response variable
df2['intercept'] = 1
observations = len(df2)
variables = df2.columns[:-1]
Xinit1 = df2.iloc[:, :-1]
#Xfin = Xc_init1.drop('Bandwidth_GB_Year', axis = 1)
y = df2['Bandwidth_GB_Year'].values
```

```
In [21]: # View number of independent variables
print ("There are", Xinit1.shape[1], "independent variables in the initial model")
```

There are 33 independent variables in the initial model.

```
In [22]: # Add y-intercept, create model, and view summary
Xc_init1 = sm.add_constant(Xinit1)
linear_regression = sm.OLS(y, Xc_init1)
fitted_model1 = linear_regression.fit()
fitted_model1.summary()
```

```
Out[22]:
```

OLS Regression Results			
<b>Dep. Variable:</b>	y	<b>R-squared:</b>	1.000
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	1.000
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	4.160e+31
<b>Date:</b>	Sun, 08 Aug 2021	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	11:48:37	<b>Log-Likelihood:</b>	2.4439e+05
<b>No. Observations:</b>	10000	<b>AIC:</b>	-4.887e+05
<b>Df Residuals:</b>	9966	<b>BIC:</b>	-4.885e+05
<b>Df Model:</b>	33		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	5.684e-13	7.78e-13	0.731	0.465	-9.56e-13	2.09e-12
<b>Population</b>	3.71e-16	4.09e-18	90.696	0.000	3.63e-16	3.79e-16

<b>Children</b>	-3.355e-13	2.91e-14	-11.543	0.000	-3.92e-13	-2.79e-13
<b>Age</b>	3.442e-15	3.03e-15	1.135	0.256	-2.5e-15	9.38e-15
<b>Income</b>	-6.245e-17	2.1e-18	-29.804	0.000	-6.66e-17	-5.83e-17
<b>Gender</b>	4.69e-13	1.2e-13	3.901	0.000	2.33e-13	7.05e-13
<b>Churn</b>	1.258e-12	1.75e-13	7.201	0.000	9.15e-13	1.6e-12
<b>Outage_sec_perweek</b>	-1.332e-15	1.99e-14	-0.067	0.947	-4.03e-14	3.76e-14
<b>Email</b>	-4.619e-14	1.95e-14	-2.365	0.018	-8.45e-14	-7.91e-15
<b>Contacts</b>	-1.092e-13	5.98e-14	-1.828	0.068	-2.26e-13	7.9e-15
<b>Yearly equip failure</b>	-4.263e-14	9.29e-14	-0.459	0.646	-2.25e-13	1.39e-13
<b>Techie</b>	1.705e-13	1.58e-13	1.076	0.282	-1.4e-13	4.81e-13
<b>Port_modem</b>	-5.755e-13	1.18e-13	-4.872	0.000	-8.07e-13	-3.44e-13
<b>Tablet</b>	-1.776e-13	1.29e-13	-1.376	0.169	-4.31e-13	7.54e-14
<b>Phone</b>	-2.203e-13	2.03e-13	-1.084	0.279	-6.19e-13	1.78e-13
<b>Multiple</b>	-3.268e-13	1.78e-13	-1.836	0.066	-6.76e-13	2.22e-14
<b>OnlineSecurity</b>	-6.75e-14	1.27e-13	-0.533	0.594	-3.16e-13	1.81e-13
<b>OnlineBackup</b>	3.553e-13	1.53e-13	2.327	0.020	5.6e-14	6.55e-13
<b>DeviceProtection</b>	8.171e-14	1.33e-13	0.616	0.538	-1.78e-13	3.42e-13
<b>TechSupport</b>	-1.421e-13	1.31e-13	-1.088	0.277	-3.98e-13	1.14e-13
<b>StreamingTV</b>	-7.248e-13	2.22e-13	-3.265	0.001	-1.16e-12	-2.9e-13
<b>StreamingMovies</b>	-2.842e-13	2.52e-13	-1.127	0.260	-7.79e-13	2.1e-13
<b>PaperlessBilling</b>	-1.137e-13	1.2e-13	-0.947	0.344	-3.49e-13	1.22e-13
<b>Tenure</b>	1.421e-13	2.55e-14	5.573	0.000	9.21e-14	1.92e-13
<b>MonthlyCharge</b>	-2.054e-15	3.82e-15	-0.537	0.591	-9.55e-15	5.44e-15
<b>Bandwidth_GB_Year</b>	1.0000	3.07e-16	3.26e+15	0.000	1.000	1.000
<b>TimelyResponse</b>	-9.948e-14	8.46e-14	-1.176	0.240	-2.65e-13	6.64e-14
<b>Fixes</b>	-1.315e-13	7.93e-14	-1.658	0.097	-2.87e-13	2.4e-14
<b>Replacements</b>	-7.461e-14	7.27e-14	-1.026	0.305	-2.17e-13	6.79e-14
<b>Reliability</b>	9.37e-14	6.5e-14	1.441	0.150	-3.37e-14	2.21e-13
<b>Options</b>	-1.066e-14	6.75e-14	-0.158	0.875	-1.43e-13	1.22e-13
<b>Respectfulness</b>	-7.105e-15	6.95e-14	-0.102	0.919	-1.43e-13	1.29e-13
<b>Courteous</b>	1.243e-14	6.58e-14	0.189	0.850	-1.16e-13	1.41e-13
<b>Listening</b>	-8.171e-14	6.25e-14	-1.306	0.191	-2.04e-13	4.09e-14
<b>Omnibus:</b>	3255.733	<b>Durbin-Watson:</b>	1.958			

<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	10848.432
<b>Skew:</b>	-1.655	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	6.884	<b>Cond. No.</b>	6.69e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Initial Multiple Linear Regression Model

5.684341886080801e-13 + 3.710139834245396e-16(Population) -3.355093980417223e-13(Children') + 3.4416913763379853e-15('Age') + 3.4416913763379853e-15('Income') + 3.4416913763379853e-15('Gender') -6.245004513516506e-17('Churn') + 4.689582056016661e-13('Outage\_sec\_perweek') + 1.2576606422953773e-12('Email') -1.3322676295501878e-15('Contacts') -4.618527782440651e-14('Yearly\_equip\_failure') -1.092459456231154e-13('Techie') -4.263256414560601e-14('Port\_modem') + 1.7053025658242404e-13('Tablet') -5.755396159656812e-13('Phone') -1.7763568394002505e-13('Multiple') -2.2026824808563106e-13('OnlineSecurity') -3.268496584496461e-13('OnlineBackup') -6.750155989720952e-14('DeviceProtection') + 3.552713678800501e-13('TechSupport') + 8.171241461241152e-14('StreamingTV') -1.4210854715202004e-13('StreamingMovies') -7.247535904753022e-13('PaperlessBilling') -2.8421709430404007e-13('Tenure') -1.1368683772161603e-13('MonthlyCharge') + 1.4210854715202004e-13('Bandwidth\_GB\_Year') -2.0539125955565396e-15 -9.947598300641403e-14('TimelyResponse') -1.3145040611561853e-13(Fixes) -7.460698725481052e-14('Replacements') + 9.370282327836321e-14 ('Reliability') -1.0658141036401503e-14('Options') -7.105427357601002e-15('Respectfulness') + 1.2434497875801753e-14('Courteous') -8.171241461241152e-14('Listening')

## D2. Model Reduction

The model has given us the highest R squared value of 1.0. However, when so many variables are present, the R squared value is inflated (Massaron & Boschetti, 2016). One sign of model instability is that the condition number is extremely high, which is a sign of multicollinearity between variables. Also, there are several variables with p-values above a 0.05 confidence interval, meaning that their coefficients are not significant to the model. I will begin by removing those variables and also the Tenure Months variable which is our response variable and therefore should not be present in the model.

This has given a new list of high p-values, and those variables will be dropped and the model re-run.

```
In [23]: Xc_init1.columns
```

```
Out[23]: Index(['const', 'Population', 'Children', 'Age', 'Income', 'Gender', 'Churn',
               'Outage_sec_perweek', 'Email', 'Contacts', 'Yearly_equip_failure',
               'Techie', 'Port_modem', 'Tablet', 'Phone', 'Multiple', 'OnlineSecurity
```

```
,
    'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
    'StreamingMovies', 'PaperlessBilling', 'Tenure', 'MonthlyCharge',
    'Bandwidth_GB_Year', 'TimelyResponse', 'Fixes', 'Replacements',
    'Reliability', 'Options', 'Respectfulness', 'Courteous', 'Listening'],
dtype='object')
```

```
In [24]: # Remove insignificant variables
drops = ['Churn', 'Tablet', 'Multiple', 'OnlineSecurity', 'TechSupport', 'StreamingTV',
        'Replacements', 'Reliability', 'Options', 'Respectfulness', 'Courteous',
        'Listening']
Xcfin = Xc_init1.drop(drops, axis = 1)
```

```
In [25]: # Re-run model
Xcfin = sm.add_constant(Xcfin)
linear_regression = sm.OLS(y, Xcfin)
fitted_model2 = linear_regression.fit()
fitted_model2.summary()
```

```
Out[25]:
```

OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	1.000
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	1.000
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	3.704e+31
<b>Date:</b>	Sun, 08 Aug 2021	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	11:48:37	<b>Log-Likelihood:</b>	2.4077e+05
<b>No. Observations:</b>	10000	<b>AIC:</b>	-4.815e+05
<b>Df Residuals:</b>	9981	<b>BIC:</b>	-4.814e+05
<b>Df Model:</b>	18		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-6.253e-12	7.65e-13	-8.168	0.000	-7.75e-12	-4.75e-12
<b>Population</b>	2.255e-16	5.87e-18	38.435	0.000	2.14e-16	2.37e-16
<b>Children</b>	-3.446e-13	3.95e-14	-8.732	0.000	-4.22e-13	-2.67e-13
<b>Age</b>	1.688e-14	4.09e-15	4.122	0.000	8.85e-15	2.49e-14
<b>Income</b>	2.819e-18	3e-18	0.939	0.348	-3.07e-18	8.71e-18
<b>Gender</b>	4.903e-13	1.7e-13	2.890	0.004	1.58e-13	8.23e-13
<b>Outage_sec_perweek</b>	7.55e-15	2.85e-14	0.265	0.791	-4.82e-14	6.33e-14
<b>Email</b>	4.619e-14	2.8e-14	1.650	0.099	-8.69e-15	1.01e-13
<b>Contacts</b>	-1.026e-13	8.57e-14	-1.197	0.231	-2.71e-13	6.54e-14
<b>Yearly equip failure</b>	-6.395e-14	1.33e-13	-0.480	0.631	-3.25e-13	1.97e-13
<b>Techie</b>	3.695e-13	2.27e-13	1.631	0.103	-7.47e-14	8.14e-13

<b>Port_modem</b>	1.847e-13	1.69e-13	1.090	0.276	-1.47e-13	5.17e-13
<b>Phone</b>	-2.558e-13	2.91e-13	-0.878	0.380	-8.27e-13	3.15e-13
<b>OnlineBackup</b>	-7.532e-13	1.79e-13	-4.215	0.000	-1.1e-12	-4.03e-13
<b>DeviceProtection</b>	-2.203e-13	1.74e-13	-1.266	0.205	-5.61e-13	1.21e-13
<b>StreamingTV</b>	-1.421e-13	1.97e-13	-0.723	0.470	-5.28e-13	2.43e-13
<b>MonthlyCharge</b>	-3.331e-15	2.4e-15	-1.385	0.166	-8.04e-15	1.38e-15
<b>Bandwidth_GB_Year</b>	1.0000	3.89e-17	2.57e+16	0.000	1.000	1.000
<b>Listening</b>	1.386e-13	8.23e-14	1.683	0.092	-2.28e-14	3e-13
<b>Omnibus:</b>	473.177	<b>Durbin-Watson:</b>	0.403			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	314.819			
<b>Skew:</b>	-0.317	<b>Prob(JB):</b>	4.34e-69			
<b>Kurtosis:</b>	2.406	<b>Cond. No.</b>	4.54e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.54e+05. This might indicate that there are

```
In [26]: # Remove insignificant variables
drops = ['Age', 'Contacts', 'Yearly_equip_failure', 'Port_modem', 'Phone', 'OnlineBackup']
Xfin2 = Xfin.drop(drops, axis = 1)
```

```
In [27]: # Re-run model
Xcfin2 = sm.add_constant(Xfin2)
linear_regression = sm.OLS(y, Xcfin2)
fitted_model2 = linear_regression.fit()
fitted_model2.summary()
```

```
Out[27]:
```

OLS Regression Results			
<b>Dep. Variable:</b>	y	<b>R-squared:</b>	1.000
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	1.000
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	9.785e+31
<b>Date:</b>	Sun, 08 Aug 2021	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	11:48:38	<b>Log-Likelihood:</b>	2.4316e+05
<b>No. Observations:</b>	10000	<b>AIC:</b>	-4.863e+05
<b>Df Residuals:</b>	9988	<b>BIC:</b>	-4.862e+05
<b>Df Model:</b>	11		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	2.16e-12	5.22e-13	4.136	0.000	1.14e-12	3.18e-12
<b>Population</b>	-3.949e-16	4.62e-18	-85.525	0.000	-4.04e-16	-3.86e-16
<b>Children</b>	1.551e-12	3.1e-14	49.962	0.000	1.49e-12	1.61e-12
<b>Income</b>	-4.12e-18	2.36e-18	-1.743	0.081	-8.75e-18	5.13e-19
<b>Gender</b>	-2.132e-14	1.33e-13	-0.160	0.873	-2.83e-13	2.4e-13
<b>Outage_sec_perweek</b>	-9.77e-15	2.24e-14	-0.436	0.663	-5.37e-14	3.41e-14
<b>Email</b>	-2.576e-14	2.2e-14	-1.169	0.242	-6.89e-14	1.74e-14
<b>Techie</b>	5.898e-13	1.78e-13	3.308	0.001	2.4e-13	9.39e-13
<b>DeviceProtection</b>	-4.334e-13	1.36e-13	-3.183	0.001	-7e-13	-1.67e-13
<b>MonthlyCharge</b>	-1.998e-15	1.58e-15	-1.268	0.205	-5.09e-15	1.09e-15
<b>Bandwidth_GB_Year</b>	1.0000	3.06e-17	3.27e+16	0.000	1.000	1.000
<b>Listening</b>	1.776e-14	6.48e-14	0.274	0.784	-1.09e-13	1.45e-13
<b>Omnibus:</b>	2256.261	<b>Durbin-Watson:</b>	1.983			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	6236.699			
<b>Skew:</b>	1.196	<b>Prob(JB):</b>	0.00			
<b>Kurtosis:</b>	6.040	<b>Cond. No.</b>	3.90e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.9e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [28]: # Remove insignificant variables
drops = ['Outage_sec_perweek', 'Email']
Xfin3 = Xfin2.drop(drops, axis = 1)
```

```
In [29]: Xcfin3 = sm.add_constant(Xfin3)
linear_regression = sm.OLS(y, Xcfin3)
fitted_model2 = linear_regression.fit()
fitted_model2.summary()
```

```
Out [29]: OLS Regression Results

Dep. Variable:          y      R-squared:      1.000
Model:          OLS      Adj. R-squared:      1.000
Method:    Least Squares      F-statistic:    6.284e+32
```

**Date:** Sun, 08 Aug 2021    **Prob (F-statistic):** 0.00  
**Time:** 11:48:38    **Log-Likelihood:** 2.5146e+05  
**No. Observations:** 10000    **AIC:** -5.029e+05  
**Df Residuals:** 9990    **BIC:** -5.028e+05  
**Df Model:** 9  
**Covariance Type:** nonrobust

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.251e-12	1.72e-13	7.266	0.000	9.13e-13	1.59e-12
<b>Population</b>	1.069e-16	2.01e-18	53.084	0.000	1.03e-16	1.11e-16
<b>Children</b>	-5.4e-13	1.35e-14	-39.881	0.000	-5.67e-13	-5.13e-13
<b>Income</b>	4.239e-17	1.03e-18	41.118	0.000	4.04e-17	4.44e-17
<b>Gender</b>	9.237e-14	5.82e-14	1.586	0.113	-2.18e-14	2.07e-13
<b>Techie</b>	-3.695e-13	7.78e-14	-4.751	0.000	-5.22e-13	-2.17e-13
<b>DeviceProtection</b>	5.365e-13	5.94e-14	9.033	0.000	4.2e-13	6.53e-13
<b>MonthlyCharge</b>	-1.443e-15	6.87e-16	-2.100	0.036	-2.79e-15	-9.61e-17
<b>Bandwidth_GB_Year</b>	1.0000	1.33e-17	7.5e+16	0.000	1.000	1.000
<b>Listening</b>	-1.315e-13	2.83e-14	-4.652	0.000	-1.87e-13	-7.61e-14
<b>Omnibus:</b>	138.923	<b>Durbin-Watson:</b>	1.381			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	155.567			
<b>Skew:</b>	-0.249	<b>Prob(JB):</b>	1.66e-34			
<b>Kurtosis:</b>	3.355	<b>Cond. No.</b>	2.96e+05			

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.96e+05. This might indicate that there are


The R value is still high. However, we have a extremely high condition number meaning multicollinearity may still be an issue. I will now look at the correlation between variables to determine if any can be removed and improve the model.

```
In [30]: # Display heatmap to view correlation between predictor variables
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(Xfin3.corr(), cmap='coolwarm')
plt.show()
```





```
In [31]: sns.heatmap(Xfin3.corr(), annot=True)
plt.show()
```



Variables are not highly correlated, which means we may have the final model. I will double check for multicollinearity using Variable Inflation Factor analysis.

In [32]:

```
# VIF dataframe
vif_data = pd.DataFrame()
vif_data["feature"] = Xfin.columns

# calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(Xfin.values, i)
                    for i in range(len(Xfin.columns))]

print(vif_data)
```

	feature	VIF
0	const	81.815779
1	Population	1.001189
2	Children	1.002558
3	Age	1.002349
4	Income	1.001485
5	Gender	1.001959
6	Outage_sec_perweek	1.001618
7	Email	1.001770
8	Contacts	1.001292
9	Yearly_equip_failure	1.001120
10	Techie	1.001485
11	Port_modem	1.001200
12	Phone	1.001595
13	OnlineBackup	1.103730
14	DeviceProtection	1.040278
15	StreamingTV	1.350033
16	MonthlyCharge	1.488423
17	Bandwidth_GB_Year	1.007482
18	Listening	1.001071

The VIF numbers for everything looks good. We have our final equation.

## D3. Reduced Multiple Regression Model

$4.093e-12 + 3.665e-17(\text{Population}) - 4.281e-13(\text{Children}) - 2.975e-14(\text{Age}) - 8.171e-13(\text{Gender}) - 1.364e-12 (\text{OnlineBackup})$

## E1. Model Comparison

The initial model of all numeric variables in the provided data set contained 33 independent variables. The inclusion of so many variables increased the R squared value, which describes how well the model fits, resulting in the perfect score of 1.0.

Variables were removed first by considering their significance, by utilizing the calculation of their p-value from the model, and by recalculating the model after each deletion to again assess the p-value. Once the insignificant variables dropped, I then examined for instances of multicollinearity. Variables that highly linearly related are problematic for predictions, which suggests that certain variables are not truly independent of each other.

To identify those variables, I tested for variance inflation factors (VIF) which measures how much individual variables are influenced by other variables (Bhandari, 2020). As with p-values, lower scores are desirable for VIF and should ideally be under 5. In the tutorials that I read, variables with similar scores were related and models improved when removing one of each related pair of variables. I did not see cases in my dataset where VIF scores were similar to be able to decide between two variables. In my final model, there were seven independent variables, meaning 60 were removed.

My final model had an R2 value of 1.0. This means that the regression model explains approximately 100.0% of the variation in the sample values of the amount of Bandwidth GB per Year.

The residuals are plotted below. The first plot, or PredictionError plot, compares the actual values from the dataset against the predicted values from by the model to see how much variance is in the model.

The line of best fit can then be compared against the 45-degree line, best identity, where the prediction exactly matches the model (McIntyre, 2018).

The difference between the two is non-existent, meaning that there is evidence that the final model is reliable. The second residuals plot shows the error of prediction.

An ideal residual plot would have points randomly distributed on either side of the axis. The plot is not random. Also, this plot includes a histogram of the error values on the right side which ideally is normally distributed around zero. The plotted histogram is nearly a perfect normal distribution curve, which suggests the higher accuracy of the model.

In [33]:

```

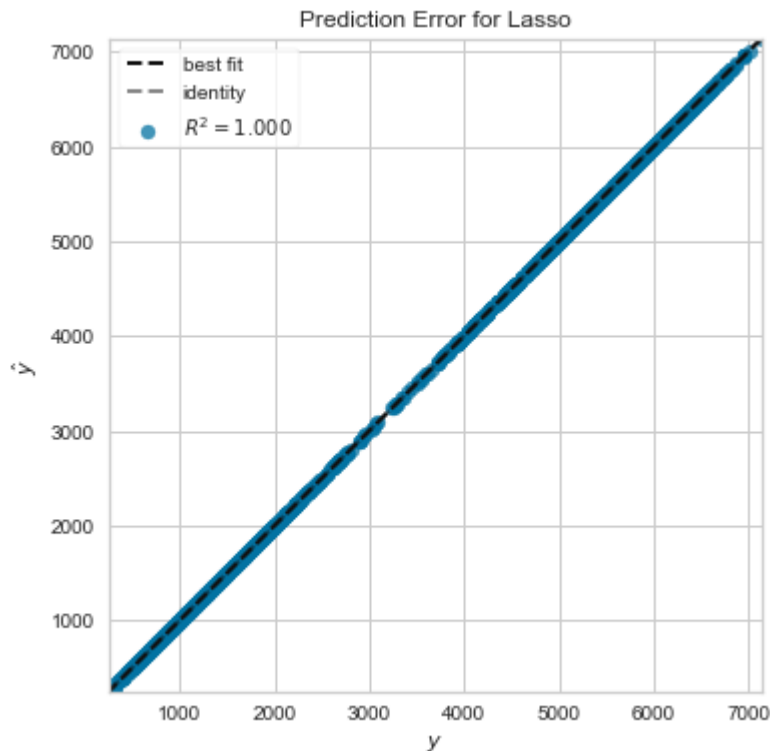
# Display prediction error plot to determine the amount of variance in model
mpl.rcParams['figure.figsize'] = (9,6)

# Create the train and test data
X_train, X_test, y_train, y_test = train_test_split(Xfin3, y, test_size=0.2,

model = Lasso()
visualizer = PredictionError(model)

visualizer.fit(X_train, y_train) # Fit the training data to the visualizer
visualizer.score(X_test, y_test) # Evaluate the model on the test data
g = visualizer.poof()           # Draw/show/poof the data

```



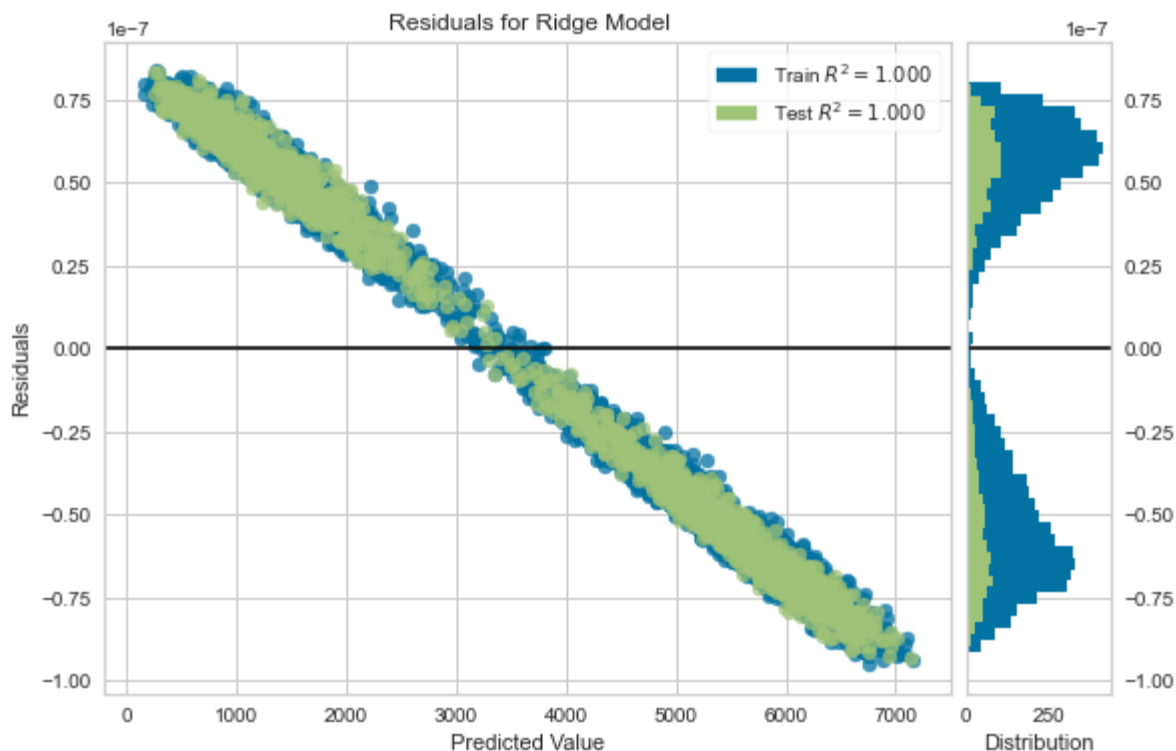
In [34]:

```

# Display residuals plot (ref 6)
model = Ridge()
visualizer = ResidualsPlot(model)

visualizer.fit(X_train, y_train) # Fit the training data to the visualizer
visualizer.score(X_test, y_test) # Evaluate the model on the test data
g = visualizer.poof()           # Draw/show/poof the data

```



## Part V: Data Summary and Implications

### F1. Results

#### Reduced Multiple Linear Regression Model

$4.093e-12 + 3.665e-17(\text{Population}) - 4.281e-13(\text{Children}) - 2.975e-14(\text{Age}) - 8.171e-13(\text{Gender}) - 1.364e-12 (\text{OnlineBackup})$

#### Results

The equation of the final regression model is:  $y = 4.093e-12 + 3.665e-17(\text{Population}) - 4.281e-13(\text{Children}) - 2.975e-14(\text{Age}) - 8.171e-13(\text{Gender}) - 1.364e-12 (\text{OnlineBackup})$

#### Interpretation of Statistical Significance of Model

The coefficients suggest that for every unit of of GB used in a year by the customer:

Population - Will increase by  $3.665e-17$  units.  
 Children - Will decrease by  $4.281e-13$  units.  
 Age- Will decrease by  $2.975e-14$  units  
 Gender- Will decrease by  $8.171e-13$  units  
 OnlineBackup - Will decrease by  $1.364e-12$  units.

#### Limitations of the Data Analysis

The data set is small, and the time period it was collected is unknown. The data set is not representative of current trends. To gain an accurate insight, the data would need to include

recent data to be train the model appropriately. The current data set is only representative of time when it was collected. There is also no time period noted for when the data was collected, so we do not know what time period is model is representative of. The data may be affected by seasonality or a bad economy, which would affect the model's relevance.

Correlation is not causation. We cannot generalize whether large amounts of gigabytes used by customers is directly related to the number of customer children, gender, or population. Each group has smaller groups within it. With age for example, we may see certain age groups using large amounts of bandwidth than other customers. Internet usage may even be different higher for some genders in areas with high population, and others with low population - population areas need to be defined and specified for a better model.

## Recommendations

Since there is a direct linear relationship between bandwidth used yearly and a customer having a online backup, I would suggest a service add on for a customer to back up their information. This could be either on the cloud or a SSD device that could be purchased with service. However, we would need to check with the customer what sorts of files/data they are storing on the cloud. The answer to that could impact what sorts of add ons we could sell or market to them.

Since children decrease the number of average GB used per year, we could suggest certain internet packages targeted towards families. For example, we could bundle kids channels in exchange for higher bandwith limits or faster speeds. This would incentivize customers to purchase larger packages that are profitable in the long run.

Since the amount of GBs used decreases with age, we could target customers with streaming packages paired with their interests, such as streaming TV services or channels that appeal to that age segment.

## Panopto Recording

<https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=137cb112-6660-48e2-960b-ad7e0114e72b>

## Sources

1. Bhandari, Aniruddha. "What Is Multicollinearity? Here's Everything You Need to Know." Analytics Vidhya, 16 Apr. 2020, [www.analyticsvidhya.com/blog/2020/03/what-is-multicollinearity/](http://www.analyticsvidhya.com/blog/2020/03/what-is-multicollinearity/).
2. Kautumn06. "Yellowbrick — Regression Visualizer Examples." Kaggle, 18 Aug. 2018, [www.kaggle.com/kautumn06/yellowbrick-regression-visualizer-examples](http://www.kaggle.com/kautumn06/yellowbrick-regression-visualizer-examples).
3. Massaron, Luca, and Alberto Boschetti. Regression Analysis with Python. Packt Publishing, 2016.
4. McKinney, W. (2018). *Python for Data Analysis*. O'Reilly.

5. Walker, Michael. *Python Data Cleaning Cookbook: Modern Techniques and Python Tools to Detect and Remove Dirty Data and Extract Key Insights*. Packt Publishing, 2020.

### Third Party Code

1. Farrell, Peter, et al. *The Statistics and Calculus with Python Workshop: A Comprehensive Introduction to Mathematics in Python for Artificial Intelligence Applications*. Packt Publishing, 2020.
2. <https://stackoverflow.com/questions/51672709/converting-no-and-yes-into-0-and-1-in-pandas-dataframe/51672855>
3. <https://analyticsvidhya.com/blog/2020/03/what-is-multicollinearity/>